

aresty_data_analysis

March 8, 2020

Intertemporal Choice: A Laboratory Investigation of Choice Behavior under Additive and Compound Wealth Growth

Barry Sopher, Department of Economics, Rutgers University

With research assistance from Himesh Buch, James Hadley and Aaron Scheiner

```
[0]: # Required libraries/packages
# =====
require(MASS)
install.packages("devtools")
library(devtools)
library(ggplot2)
```

```
[0]: main_data <- read.csv(file="/content/stacked-data-ergodicity.csv",header = TRUE)
summary(main_data)
```

first_name	last_name	login_id	start_time
Joseph : 42	Patel : 63	ac1887 : 21	2/26/2020 12:13:294
Nicholas: 42	Shah : 63	acm298 : 21	2/25/2020 10:16:168
Ryan : 42	Parekh : 42	ah1123 : 21	2/26/2020 10:13:147
Aditya : 21	Ackerley: 21	aka99 : 21	2/24/2020 10:13:126
Akil : 21	Ahmed : 21	al1104 : 21	2/24/2020 12:10:126
Alan : 21	Ahmmmed : 21	aman.s.gupta: 21	2/25/2020 12:13:105
(Other) :1365	(Other) :1323	(Other) :1428	(Other) :588

end_time	duration	treatment	session
2/24/2020 10:26: 63	Min. : 3.00	Min. :1.000	Min. :1.000
2/26/2020 12:27: 63	1st Qu.: 8.00	1st Qu.:1.000	1st Qu.:2.000
2/24/2020 10:23: 42	Median :12.50	Median :1.000	Median :4.000
2/24/2020 12:18: 42	Mean :12.99	Mean :1.473	Mean :3.676
2/25/2020 10:25: 42	3rd Qu.:16.00	3rd Qu.:2.000	3rd Qu.:5.000
2/25/2020 12:21: 42	Max. :34.00	Max. :2.000	Max. :6.000
(Other) :1260			

subject	question
Min. : 1.0	Min. : 1.00
1st Qu.:19.0	1st Qu.: 6.00
Median :37.5	Median :11.00
Mean :37.5	Mean :10.98
3rd Qu.:56.0	3rd Qu.:16.00

Max. :74.0 Max. :21.00

A 50% chance of Option A and a 50% chance of Option B.

Option A: \$2.00 every 3 days. A 3 percent rate of interest on accumulated earnings every time

Option A: \$0.29 every 1 days. Accumulations end in 70 days.

Option A: \$0.67 every 1 days. Accumulations end in 30 days.

Option B: \$7.00 every 9 days. Accumulations end in 30 days.

Option B: \$3.00 every 7 days. Accumulations end in 50 days.

(Other)

choice	Xa	Xb	H	D
Min. :0.000	Min. :0.290	Min. :0.430	Min. :1.000	Min. :2
1st Qu.:0.000	1st Qu.:1.330	1st Qu.:2.000	1st Qu.:3.000	1st Qu.:2
Median :1.000	Median :2.000	Median :3.000	Median :5.000	Median :2
Mean :0.546	Mean :2.248	Mean :3.349	Mean :4.946	Mean :2
3rd Qu.:1.000	3rd Qu.:2.800	3rd Qu.:4.200	3rd Qu.:7.000	3rd Qu.:2
Max. :1.000	Max. :6.000	Max. :9.000	Max. :9.000	Max. :2

T	cumXa	cumXb	idealChoice
Min. :30.00	Min. :16.48	Min. : 9.60	A:709
1st Qu.:30.00	1st Qu.:19.52	1st Qu.:18.00	B:845
Median :50.00	Median :20.00	Median :21.00	
Mean :49.55	Mean :20.58	Mean :19.82	
3rd Qu.:70.00	3rd Qu.:21.34	3rd Qu.:21.97	
Max. :70.00	Max. :27.82	Max. :26.27	

Looking at the above code snippet and its result, we will only focus on our response variable (choice) and all independent variables (Xa, Xb, H, D, T) in this course of action. These results might be useful in future

```
[0]: # Regression
# =====
lmod <- lm(choice ~ Xa+Xb+H+D+T , data=main_data) # linear model
# regression
summary(lmod)
```

Call:

```
lm(formula = choice ~ Xa + Xb + H + D + T, data = main_data)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.00700	-0.31518	0.04834	0.33846	0.87492

Coefficients: (1 not defined because of singularities)

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.128372	0.075703	1.696	0.0901 .

```

Xa      -0.182942    0.116654   -1.568    0.1170
Xb       0.091822    0.073814    1.244    0.2137
H       0.128112    0.012905    9.927   <2e-16 ***
D              NA              NA              NA              NA
T      -0.002266    0.001415   -1.602    0.1094
---

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4186 on 1549 degrees of freedom

Multiple R-squared: 0.2533, Adjusted R-squared: 0.2514

F-statistic: 131.4 on 4 and 1549 DF, p-value: < 2.2e-16

Results above show that, among all of our independent variables, Horizon (H) has the least P-Value, making it the most important variable which affects the response variable the most. We can even ignore the other variables

```
[0]: anova(lmod)
```

		Df <int>	Sum Sq <dbl>	Mean Sq <dbl>	F value <dbl>	Pr(>F) <dbl>
A anova: 5 × 5	Xa	1	48.816593	48.8165934	278.637384	1.187734e-57
	Xb	1	2.291036	2.2910356	13.076868	3.085475e-04
	H	1	40.521987	40.5219869	231.293084	8.655774e-49
	T	1	0.449608	0.4496080	2.566292	1.093673e-01
	Residuals	1549	271.381040	0.1751976	NA	NA

```
[0]: coef(lmod)
```

```

(Intercept) 0.128372093392043 Xa -0.182941841781605 Xb 0.0918220441099204 H
0.128112081222266 D <NA> T -0.00226645355776882

```

The above result help us come up with the multiple linear regression model

```

[0]: confint(lmod)           # CI at 95%
confint(lmod, level=.99)    # CI at 99%


```


		2.5 %	97.5 %
A matrix: 6 × 2 of type dbl	(Intercept)	-0.020118423	0.2768626099
	Xa	-0.411757689	0.0458740057
	Xb	-0.052963646	0.2366077338
	H	0.102798400	0.1534257627
	D	NA	NA
	T	-0.005041571	0.0005086641


	0.5 %	99.5 %
(Intercept)	-0.066865600	0.323609787
Xa	-0.483792559	0.117908875
Xb	-0.098544485	0.282188573
H	0.094829250	0.161394912
D	NA	NA
T	-0.005915222	0.001382315


We have computed the 95% and 99% confidence interval for our model. One can confirm the results by comparing the P-Values. This result can be very useful for the hypothesis testing, as it will help to determine whether to ignore the null hypothesis or not

```
[0]: plot(lmod)
```









These are the graphs generated by the regression analysis. Each graph helps us understand different things about residuals and distribution of the data. Here is what each graph represents:

1. **Residual(errors) vs Fitted:** linearity of the model and shows if residuals have non-linear patterns
2. **Normal Q-Q plot (Normal Quantile-Quantile plot):** shows if residuals are normally distributed
3. **Scale-Location:** shows if residuals are spread equally along the ranges of predictors
4. **Residuals vs Leverage:** helps us to find influential cases (robust regression can also help)

```
[0]: lmod_g <- glm(choice ~ Xa+Xb+H+D+T , data=main_data) # generalized
      ↪ linear model
      summary(lmod_g)
```

```

Call:
glm(formula = choice ~ Xa + Xb + H + D + T, data = main_data)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.00700  -0.31518   0.04834   0.33846   0.87492

Coefficients: (1 not defined because of singularities)
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.128372   0.075703   1.696   0.0901 .
Xa          -0.182942   0.116654  -1.568   0.1170
Xb           0.091822   0.073814   1.244   0.2137
H            0.128112   0.012905   9.927  <2e-16 ***
D              NA         NA         NA      NA
T           -0.002266   0.001415  -1.602   0.1094
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 0.1751976)

Null deviance: 363.46  on 1553  degrees of freedom
Residual deviance: 271.38  on 1549  degrees of freedom
AIC: 1710.2

Number of Fisher Scoring iterations: 2

```

Here is a generalized linear modal. The modal lmod was a simple linear modal, where this modal helps us understand the exponentiality of the distribution

```
[0]: anova(lmod_g)
```

		Df	Deviance	Resid. Df	Resid. Dev
		<int>	<dbl>	<int>	<dbl>
A anova: 6 × 4	NULL	NA	NA	1553	363.4603
	Xa	1	48.816593	1552	314.6437
	Xb	1	2.291036	1551	312.3526
	H	1	40.521987	1550	271.8306
	D	0	0.000000	1550	271.8306
	T	1	0.449608	1549	271.3810

```
[0]: confint(lmod_g)           # CI at 95%
      confint(lmod_g, level=.99) # CI at 99%
```

Waiting for profiling to be done...

		2.5 %	97.5 %
A matrix: 6 × 2 of type dbl	(Intercept)	-0.020002396	0.2767465833
	Xa	-0.411578898	0.0456952149
	Xb	-0.052850514	0.2364946020
	H	0.102818179	0.1534059832
	D	NA	NA
	T	-0.005039403	0.0005064957

Waiting for profiling to be done...

		0.5 %	99.5 %
A matrix: 6 × 2 of type dbl	(Intercept)	-0.066625040	0.32336923
	Xa	-0.483421869	0.11753818
	Xb	-0.098309926	0.28195401
	H	0.094870259	0.16135390
	D	NA	NA
	T	-0.005910727	0.00137782

In the above snippets, we looked at generalized linear modal and its anova along with the confidence intervals

```
[0]: lmod_gnb <- glm.nb(choice ~ Xa+Xb+H+D+T , data=main_data) # negative binomial
      ↪ model
```

```
[0]: summary(lmod_gnb)
```

Call:

```
glm.nb(formula = choice ~ Xa + Xb + H + D + T, data = main_data,
       init.theta = 20592.51732, link = log)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.52843	-0.82247	0.05077	0.46607	1.20075

Coefficients: (1 not defined because of singularities)

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.450862	0.281308	-5.158	2.5e-07 ***
Xa	-0.514161	0.526948	-0.976	0.329
Xb	0.285997	0.339990	0.841	0.400
H	0.240630	0.041925	5.740	9.5e-09 ***
D	NA	NA	NA	NA
T	-0.005057	0.005276	-0.958	0.338

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for Negative Binomial(20592.52) family taken to be 1)

Null deviance: 1026.87 on 1553 degrees of freedom

Residual deviance: 855.72 on 1549 degrees of freedom
AIC: 2543.7

Number of Fisher Scoring iterations: 1

Theta: 20593
Std. Err.: 68329
Warning while fitting theta: iteration limit reached

2 x log-likelihood: -2531.744

```
[0]: anova(lmod_gnb)
```

Warning message in anova.negbin(lmod_gnb):
"tests made without re-estimating 'theta'"

		Df	Deviance	Resid. Df	Resid. Dev	Pr(>Chi)
		<int>	<dbl>	<int>	<dbl>	<dbl>
A anova: 6 × 5	NULL	NA	NA	1553	1026.8667	NA
	Xa	1	82.2917733	1552	944.5749	1.174098e-19
	Xb	1	4.8276907	1551	939.7472	2.800612e-02
	H	1	83.1065249	1550	856.6407	7.774863e-20
	D	0	0.0000000	1550	856.6407	NA
	T	1	0.9220767	1549	855.7186	3.369303e-01

```
[0]: confint(lmod_gnb)           # CI at 95%
     confint(lmod_gnb, level=.99) # CI at 99%
```

Waiting for profiling to be done...

		2.5 %	97.5 %
A matrix: 6 × 2 of type dbl	(Intercept)	-2.0067405	-0.903729261
	Xa	-1.6512186	0.437843122
	Xb	-0.3220629	1.026101982
	H	0.1589407	0.323315416
	D	NA	NA
	T	-0.0154445	0.005246005

Waiting for profiling to be done...

		0.5 %	99.5 %
A matrix: 6 × 2 of type dbl	(Intercept)	-2.18335716	-0.733467723
	Xa	-2.05922968	0.708823322
	Xb	-0.49258754	1.294168864
	H	0.13346256	0.349516468
	D	NA	NA
	T	-0.01872919	0.008469432

In the above snippets, we also looked at negative binomial distribution model and its anova along with the confidence intervals

Now we are considering different types of regression modals

```
[0]: install.packages("olsrr")

[0]: # Different types of regression
# =====
require(olsrr)
main_data <- read.csv(file="/stacked-data-ergodicity.csv",header = TRUE)
lmod <- lm(choice ~ Xa+Xb+H+D+T , data=main_data) # linear model ↵
↵ regression
k <- ols_step_all_possible(lmod) # all possible regressions
```

Here's the output generated by R for k in the above snippet:

	Index	N	Predictors	R-Square	Adj. R-Square	Mallow's Cp
3	1	1	H	0.251406437	0.2509240962	3.012463
2	2	1	Xb	0.137924879	0.1373694180	238.438312
1	3	1	Xa	0.134310675	0.1337528854	245.936244
5	4	1	T	0.000129974	-0.0005142721	524.303987
4	5	1	D	0.000000000	0.0000000000	522.573628
7	6	2	Xa H	0.251644988	0.2506799914	4.517572
10	7	2	Xb H	0.251573823	0.2506087346	4.665209
14	8	2	H T	0.251406636	0.2504413317	5.012051
13	9	2	H D	0.251406437	0.2509240962	3.012463
9	10	2	Xa T	0.205578190	0.2045537901	100.086537
12	11	2	Xb T	0.205295974	0.2042712108	100.672014
6	12	2	Xa Xb	0.140614075	0.1395059054	234.859375
11	13	2	Xb D	0.137924879	0.1373694180	238.438312
8	14	2	Xa D	0.134310675	0.1337528854	245.936244
15	15	2	D T	0.000129974	-0.0005142721	524.303987
20	16	3	Xa H T	0.252594638	0.2511480467	4.547454
23	17	3	Xb H T	0.252155056	0.2507076139	5.459399
16	18	3	Xa Xb H	0.252103531	0.2506559893	5.566292
19	19	3	Xa H D	0.251644988	0.2506799914	4.517572
22	20	3	Xb H D	0.251573823	0.2506087346	4.665209
25	21	3	H D T	0.251406636	0.2504413317	5.012051
18	22	3	Xa Xb T	0.205838121	0.2043010331	101.547291
21	23	3	Xa D T	0.205578190	0.2045537901	100.086537


24	24	3		Xb	D	T	0.205295974	0.2042712108	100.672014
17	25	3		Xa	Xb	D	0.140614075	0.1395059054	234.859375
27	26	4		Xa	Xb	H	T	0.253340552	0.2514124450
29	27	4		Xa	H	D	T	0.252594638	0.2511480467
30	28	4		Xb	H	D	T	0.252155056	0.2507076139
26	29	4		Xa	Xb	H	D	0.252103531	0.2506559893
28	30	4		Xa	Xb	D	T	0.205838121	0.2043010331
31	31	5		Xa	Xb	H	D	T	0.253340552
									0.2514124450
									5.000000

```
[24]: k <- ols_step_best_subset(lmod, print_plot = TRUE) # best subset
      ↪regression
      plot(k)
```


```
[[1]]
NULL
```

```
[[2]]
NULL
```


R-Square




C(p)




Adj. R-Square




AIC



SBIC



SBC



Here is the output generated by R for k in the above snippet:

Best Subsets Regression

Model Index	Predictors
1	H
2	Xa H
3	Xa H T
4	Xa Xb H T
5	Xa Xb H D T

Subsets Regression Summary

Model	R-Square	Adj. R-Square	Pred R-Square	C(p)	AIC	SBIC
1	0.2514	0.2509	0.2496	3.0125	1708.2523	NA
2	0.2516	0.2507	0.2487	4.5176	1709.7570	NA
3	0.2526	0.2511	0.2488	4.5475	1709.7837	NA
4	0.2533	0.2514	0.2486	5.0000	1710.2321	NA
5	0.2533	0.2514	0.2486	5.0000	1712.2321	NA

Model	SBC	MSEP	FPE	HSP	APC
1	1724.2980	272.4346	0.1755	1e-04	0.7505
2	1731.1513	272.5235	0.1757	1e-04	0.7513
3	1736.5267	272.3534	0.1757	1e-04	0.7513
4	1742.3236	272.2574	0.1758	1e-04	0.7515
5	1749.6722	272.2574	0.1758	1e-04	0.7515

AIC: Akaike Information Criteria

SBIC: Sawa's Bayesian Information Criteria

SBC: Schwarz Bayesian Criteria

MSEP: Estimated error of prediction, assuming multivariate normality

FPE: Final Prediction Error

HSP: Hocking's Sp

APC: Amemiya Prediction Criteria

```
[26]: # forward stepwise regression
# =====
k <- ols_step_forward_p(lmod,details=TRUE)
k
```

Forward Selection Method

Candidate Terms:

1. Xa
2. Xb
3. H
4. D
5. T

We are selecting variables based on p value...

Note: model has aliased coefficients

sums of squares computed by model comparison

Forward Selection: Step 1

- H

Model Summary

R	0.501	RMSE	0.419
R-Squared	0.251	Coef. Var	76.684
Adj. R-Squared	0.251	MSE	0.175
Pred R-Squared	0.250	MAE	0.358

RMSE: Root Mean Square Error

MSE: Mean Square Error

MAE: Mean Absolute Error

ANOVA

	Sum of Squares	DF	Mean Square	F	Sig.
Regression	91.376	1	91.376	521.221	0.0000
Residual	272.084	1552	0.175		
Total	363.460	1553			

Parameter Estimates

model	Beta	Std. Error	Std. Beta	t	Sig	lower upper
(Intercept)	0.009	0.026		0.368	0.713	-0.041 0.060
H	0.108	0.005	0.501	22.830	0.000	0.099 0.118

Note: model has aliased coefficients

sums of squares computed by model comparison

No more variables to be added.

Variables Entered:

+ H

Final Model Output

Model Summary

R	0.501	RMSE	0.419
R-Squared	0.251	Coef. Var	76.684
Adj. R-Squared	0.251	MSE	0.175
Pred R-Squared	0.250	MAE	0.358

RMSE: Root Mean Square Error
MSE: Mean Square Error
MAE: Mean Absolute Error

ANOVA

	Sum of Squares	DF	Mean Square	F	Sig.
Regression	91.376	1	91.376	521.221	0.0000
Residual	272.084	1552	0.175		
Total	363.460	1553			

Parameter Estimates

model	Beta	Std. Error	Std. Beta	t	Sig	lower upper
(Intercept)	0.009	0.026		0.368	0.713	-0.041 0.060
H	0.108	0.005	0.501	22.830	0.000	0.099 0.118

Selection Summary

Step	Variable Entered	R-Square	Adj. R-Square	C(p)	AIC	RMSE
1	H	0.2514	0.2509	3.0125	1708.2523	0.4187

[27]: `plot(k)`

geom_path: Each group consists of only one observation. Do you need to adjust the group aesthetic?

geom_path: Each group consists of only one observation. Do you need to adjust the group aesthetic?

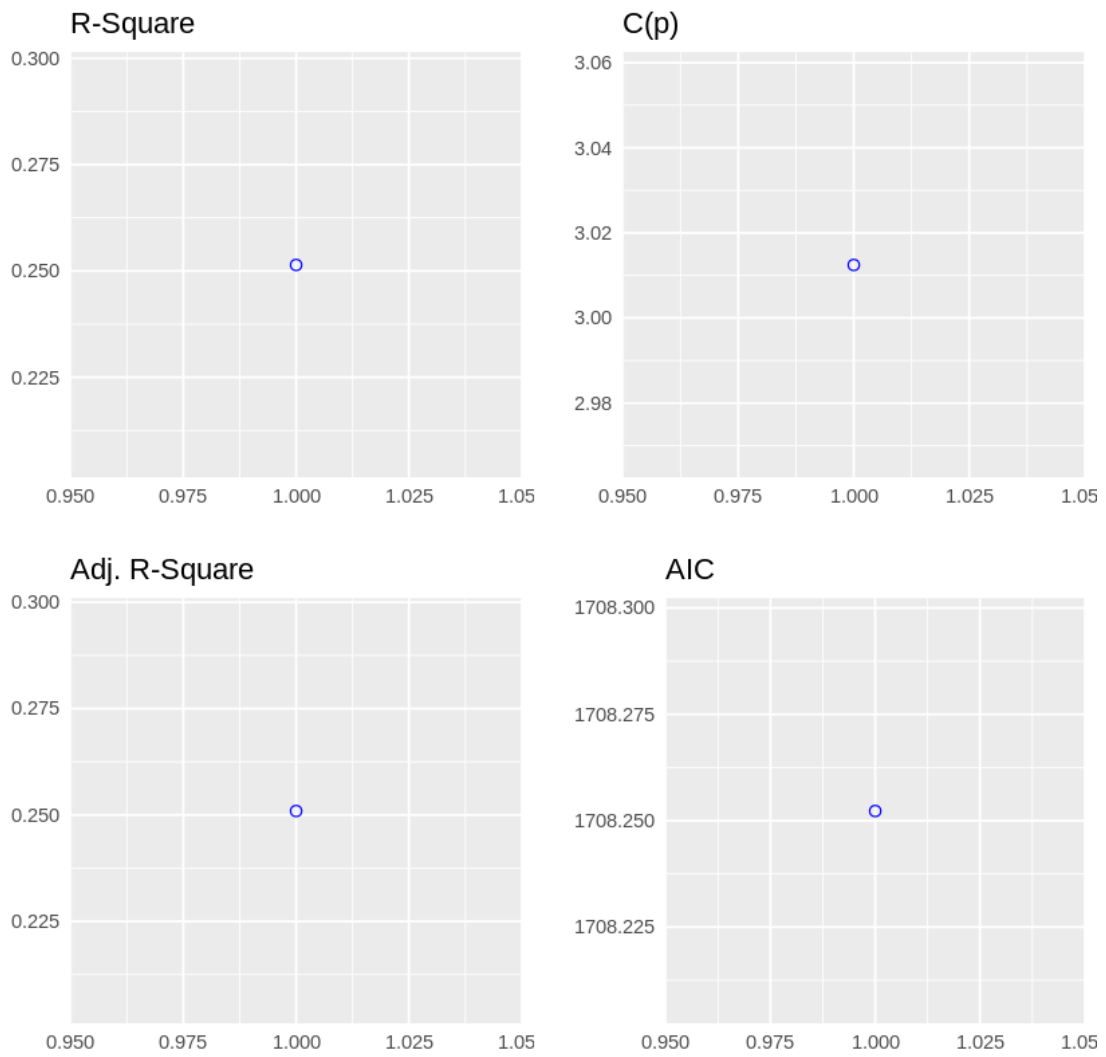
geom_path: Each group consists of only one observation. Do you need to adjust the group aesthetic?

geom_path: Each group consists of only one observation. Do you need to adjust the group aesthetic?


geom_path: Each group consists of only one observation. Do you need to adjust the group aesthetic?

[[1]]
NULL


[[2]]
NULL



SBIC



SBC



Explanation: Build regression model from a set of candidate predictor variables by entering predictors based on p values, in a stepwise manner until there is no variable left to enter any more. The model should include all the candidate predictor variables. If `details` is set to `TRUE`, each step is displayed.

```
[28]: # backward stepwise regression
# =====
k <- ols_step_backward_p(lmod,details=TRUE)
k
```

Backward Elimination Method

Candidate Terms:

1 . Xa
 2 . Xb
 3 . H
 4 . D
 5 . T

We are eliminating variables based on p value...

Note: model has aliased coefficients
 sums of squares computed by model comparison

No more variables satisfy the condition of p value = 0.3

Variables Removed:

Final Model Output

Model Summary			
R	0.503	RMSE	0.419
R-Squared	0.253	Coef. Var	76.659
Adj. R-Squared	0.251	MSE	0.175
Pred R-Squared	0.249	MAE	0.357

RMSE: Root Mean Square Error

MSE: Mean Square Error

MAE: Mean Absolute Error

ANOVA					
	Sum of Squares	DF	Mean Square	F	Sig.
Regression	92.079	4	23.020	131.393	0.0000
Residual	271.381	1549	0.175		
Total	363.460	1553			

Parameter Estimates	

model	Beta	Std. Error	Std. Beta	t	Sig	lower	upper

(Intercept)	0.128	0.076		1.696	0.090	-0.020	0.277
Xa	-0.183	0.117	-0.499	-1.568	0.117	-0.412	0.046
Xb	0.092	0.074	0.378	1.244	0.214	-0.053	0.237
H	0.128	0.013	0.592	9.927	0.000	0.103	0.153
D	NA	0.001	0.000	-1.602	0.109	NA	NA
T	-0.002	NA	-6.278	NA	NA	-0.005	0.001


```
[1] "No variables have been removed from the model."
```

Explanation: Build regression model from a set of candidate predictor variables by removing predictors based on p values, in a stepwise manner until there is no variable left to remove any more. The model should include all the candidate predictor variables. If details is set to TRUE, each step is displayed

```
[30]: # sequential (both directions) stepwise regression
# =====
k <- ols_step_both_p(lmod,details=TRUE)
k
```

Stepwise Selection Method

Candidate Terms:

1. Xa
2. Xb
3. H
4. D
5. T

We are selecting variables based on p value...

Note: model has aliased coefficients
 sums of squares computed by model comparison

Stepwise Selection: Step 1

- H added

Model Summary

R	0.501	RMSE	0.419
R-Squared	0.251	Coef. Var	76.684
Adj. R-Squared	0.251	MSE	0.175
Pred R-Squared	0.250	MAE	0.358

RMSE: Root Mean Square Error

MSE: Mean Square Error

MAE: Mean Absolute Error

ANOVA

	Sum of Squares	DF	Mean Square	F	Sig.
Regression	91.376	1	91.376	521.221	0.0000
Residual	272.084	1552	0.175		
Total	363.460	1553			

Parameter Estimates

model	Beta	Std. Error	Std. Beta	t	Sig.	lower upper
(Intercept)	0.009	0.026		0.368	0.713	-0.041 0.060
H	0.108	0.005	0.501	22.830	0.000	0.099 0.118

Note: model has aliased coefficients

sums of squares computed by model comparison

No more variables to be added/removed.

Final Model Output

Model Summary

R	0.501	RMSE	0.419
R-Squared	0.251	Coef. Var	76.684
Adj. R-Squared	0.251	MSE	0.175
Pred R-Squared	0.250	MAE	0.358

RMSE: Root Mean Square Error
MSE: Mean Square Error
MAE: Mean Absolute Error

ANOVA

	Sum of Squares	DF	Mean Square	F	Sig.
Regression	91.376	1	91.376	521.221	0.0000
Residual	272.084	1552	0.175		
Total	363.460	1553			

Parameter Estimates

model	Beta	Std. Error	Std. Beta	t	Sig	lower upper
(Intercept)	0.009	0.026		0.368	0.713	-0.041 0.060
H	0.108	0.005	0.501	22.830	0.000	0.099 0.118

Stepwise Selection Summary

Step	Variable	Added/ Removed	R-Square	Adj. R-Square	C(p)	AIC	RMSE
1	H	addition	0.251	0.251	3.0120	1708.2523	0.4187

```
[31]: plot(k)
```

```
geom_path: Each group consists of only one observation. Do you need to adjust  
the group aesthetic?
```

```
geom_path: Each group consists of only one observation. Do you need to adjust  
the group aesthetic?
```

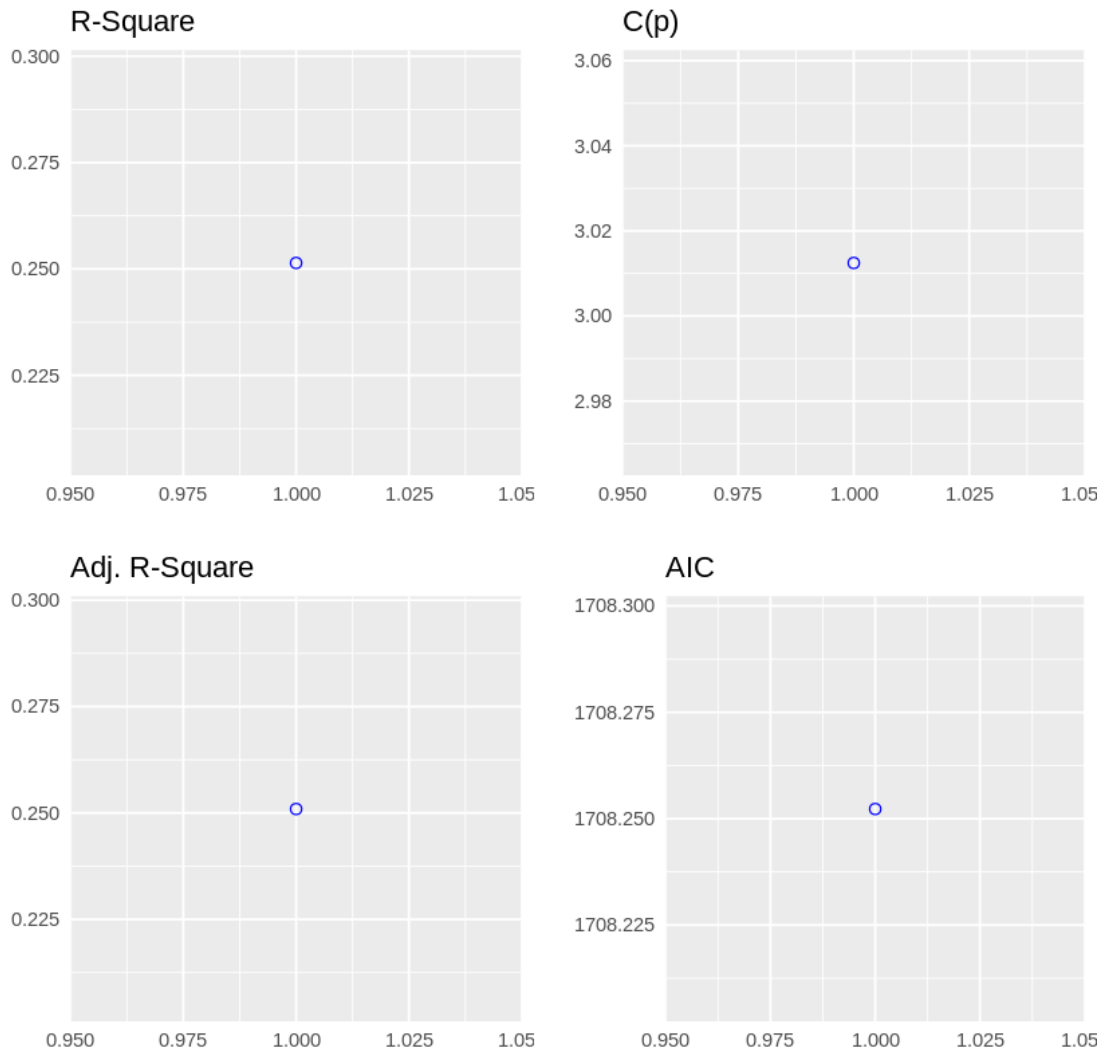
```
geom_path: Each group consists of only one observation. Do you need to adjust  
the group aesthetic?
```

```
geom_path: Each group consists of only one observation. Do you need to adjust  
the group aesthetic?
```


```
geom_path: Each group consists of only one observation. Do you need to adjust  
the group aesthetic?
```

```
[[1]]  
NULL
```


```
[[2]]  
NULL
```



SBIC



SBC



Explanation: Build regression model from a set of candidate predictor variables by entering and removing predictors based on p values, in a stepwise manner until there is no variable left to enter or remove any more. The model should include all the candidate predictor variables. If details is set to `TRUE`, each step is displayed.

```
[32]: # Stepwise models based on Akaike Information Criteria criterion
# =====
k <- ols_step_forward_aic(lmod)      # forward selection
k
```

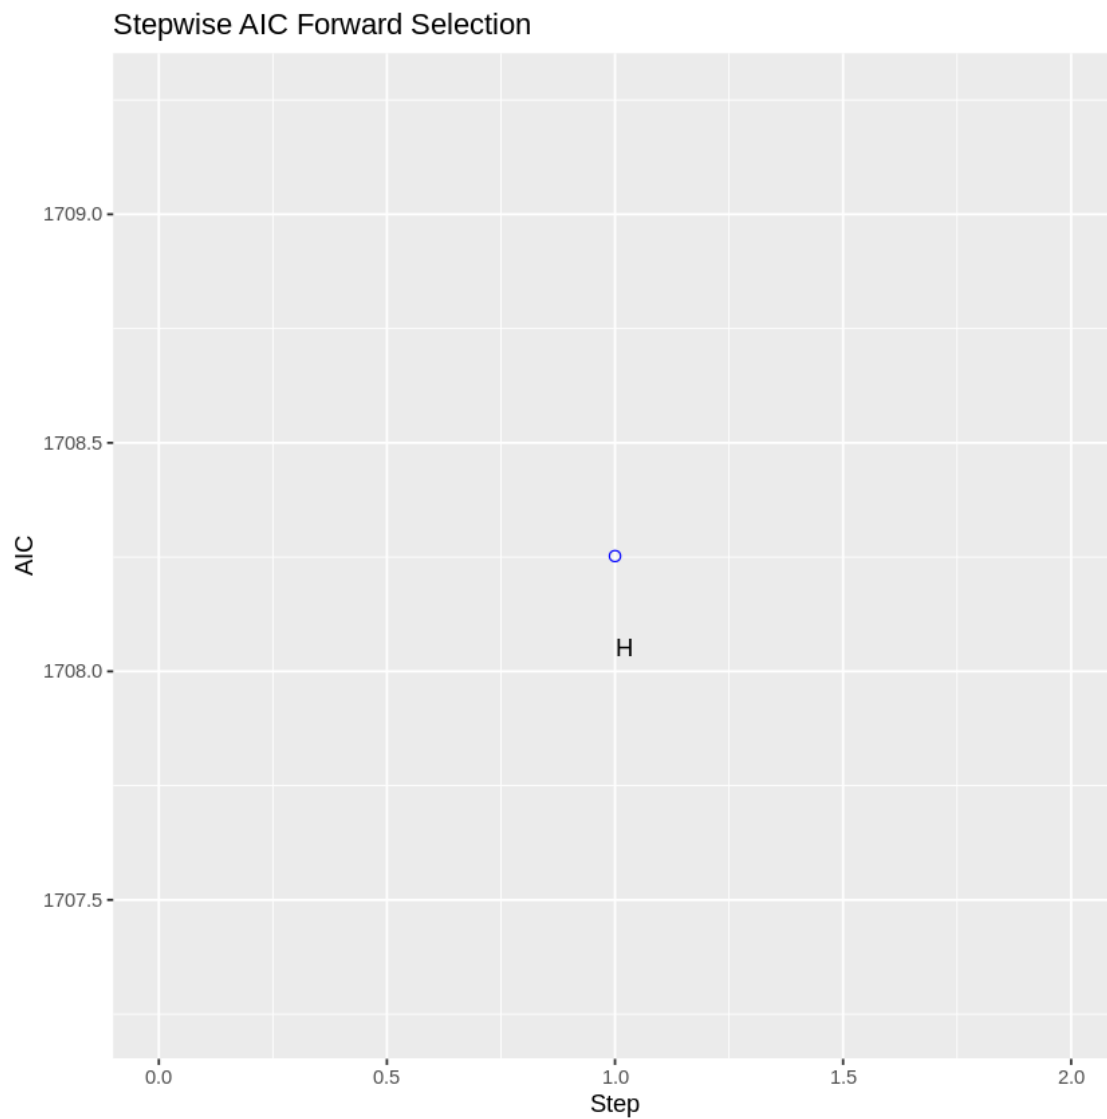
Selection Summary

Variable	AIC	Sum Sq	RSS	R-Sq	Adj. R-Sq
----------	-----	--------	-----	------	-----------

H	1708.252	91.376	272.084	0.25141	0.25092

```
[33]: plot(k)
```

geom_path: Each group consists of only one observation. Do you need to adjust the group aesthetic?



```
[35]: k <- ols_step_both_aic(lmod)      # both sides
      k
```

Stepwise Summary						
Variable	Method	AIC	RSS	Sum Sq	R-Sq	Adj. R-Sq
H	addition	1708.252	272.084	91.376	0.25141	0.25092

```
[37]: plot(k)
```

geom_path: Each group consists of only one observation. Do you need to adjust the group aesthetic?

