

# Linear regression by gradient descent

July 26, 2012

By [Christopher Bare](#)



[This article was first published on [Digithead's Lab Notebook](#), and kindly contributed to [R-bloggers](#). (You can report issue about the content on this page [here](#))

Want to share your content on R-bloggers? [click here](#) if you have a blog, or [here](#) if you don't.

Share

Tweet

In [Andrew Ng's Machine Learning class](#), the first section demonstrates [gradient descent](#) by using it on a familiar problem, that of fitting a linear function to data.

Let's start off, by generating some bogus data with known characteristics. Let's make y just a noisy version of x. Let's also add 3 to give the intercept term something to do.

```
# generate random data in which y is a noisy function of x
x <- runif(1000, -5, 5)
y <- x + rnorm(1000) + 3
```


```
# fit a linear model
res <- lm( y ~ x )
print(res)
```

```
Call:
lm(formula = y ~ x)
```

```
Coefficients:
(Intercept)          x
    2.9930      0.9981
```

Fitting a linear model, we should get a slope of 1 and an intercept of 3. Sure enough, we get pretty close. Let's plot it and see how it looks.

```
# plot the data and the model
plot(x,y, col=rgb(0.2,0.4,0.6,0.4), main='Linear regression by gradient descent')
abline(res, col='blue')
```



As a learning exercise, we'll do the same thing using gradient descent. As [discussed previously](#), the main idea is to take the partial derivative of the cost function with respect to theta. That gradient, multiplied by a learning rate, becomes the update rule for the estimated values of the parameters. Iterate and things should converge nicely.

```
# squared error cost function
cost <- function(X, y, theta) {
  sum( (X %*% theta - y)^2 ) / (2*length(y))
}

# Learning rate and iteration limit
alpha <- 0.01
num_iters <- 1000

# keep history
cost_history <- double(num_iters)
theta_history <- list(num_iters)

# initialize coefficients
theta <- matrix(c(0,0), nrow=2)

# add a column of 1's for the intercept coefficient
X <- cbind(1, matrix(x))

# gradient descent
for (i in 1:num_iters) {
  error <- (X %*% theta - y)
  delta <- t(X) %*% error / length(y)
  theta <- theta - alpha * delta
  cost_history[i] <- cost(X, y, theta)
  theta_history[[i]] <- theta
}


print(theta)

      [,1]
[1,] 2.9928978
[2,] 0.9981226
```

As expected, theta winds up with the same values as [lm](#) returned. Let's do some more plotting:


```
# plot data and converging fit
plot(x,y, col=rgb(0.2,0.4,0.6,0.4), main='Linear regression by gradient descent')
```

```
for (i in c(1,3,6,10,14,seq(20,num_iters,by=10))) {
  abline(coef=theta_history[[i]], col=rgb(0.8,0,0,0.3))
}
abline(coef=theta, col='blue')
```



Taking a look at how quickly the cost decreases, I might have done with fewer iterations.

```
plot(cost_history, type='line', col='blue', lwd=2, main='Cost function', ylab='cost', xlab='Iterations')
```



That was easy enough. The next step is to look into some of the more advanced [optimization methods](#) available within R. I'll try to translate more of the [Machine Learning class](#) into R. I know [others](#) are doing that as well.